

question4.py

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4
5 # Load data from the provided CSV file.
6 data = pd.read_csv("HW2_linear_data.csv")
7 # Load the first (input/feature) column into the vector X.
8 X = data.iloc[:, 0].values
9 # Load the second (output/target) column into the vector Y.
10 Y = data.iloc[:, 1].values
11
12 # Initialize the parameters to 0.
13 # The m value is the slope (i.e., weight), and the c value
14 # is the y-intercept (i.e., bias).
15 m = 0
16 c = 0
17
18 # Set learning rate and number of epochs as described in the problem statement.
19 learning_rate = 0.0001
20 epochs = 1000
21
22 # Count the number of input values.
23 n = len(X)
24
25 # Calculate the gradient descent.
26 for e in range(epochs):
27     # Calculate the predicted value
28     Y_pred = m * X + c
29     # Calculate the error relative to ground truth.
30     error = Y_pred - Y
31     # Use the error value to calculate the MSE.
32     mse = np.mean(error**2)
33
34     # Update the slope/weight and intercept/bias based on the
35     # learning rate and the error.
36
37     # When updating m, the term "np.sum(error * X)" represents
38     # the gradient of the MSE with respect to m.
39     m -= learning_rate * (2/n) * np.sum(error * X) # Update slope
40     # When updating c, the term "np.sum(error)" represents the
41     # gradient of the MSE with respect to c.
42     c -= learning_rate * (2/n) * np.sum(error) # Update intercept
43
44     # Print the MSE value every 100 epochs to see changes.
45     # The MSE at epoch 0 should be very large because of the initial
46     # values of 0 for both m and c, with a sharp drop at epoch
47     # epoch 100, followed by small incremental improvements for the
```

```
48     # remaining epochs.
49     if e % 100 == 0:
50         print(f"Epoch {e}: MSE = {mse}")
51
52 # Generate the final predictions using the trained values of m and c.
53 Y_pred = m * X + c
54
55 # Plot the results.
56 # Generate a scatter plot of the ground-truth data.
57 plt.scatter(X, Y, color="blue", label="Actual Data")
58 # Plot the predictions as a line.
59 plt.plot(X, Y_pred, color="red", label="Fitted Line")
60 plt.xlabel("X")
61 plt.ylabel("Y")
62 plt.title("Linear Regression")
63 plt.show()
64
65 # Print the final trained parameters.
66 print(f"Final slope (m): {m}")
67 print(f"Final intercept (c): {c}")
68
```